



Efficiency comparison — proposed Custom V2 (Hub v2) vs nine baselines

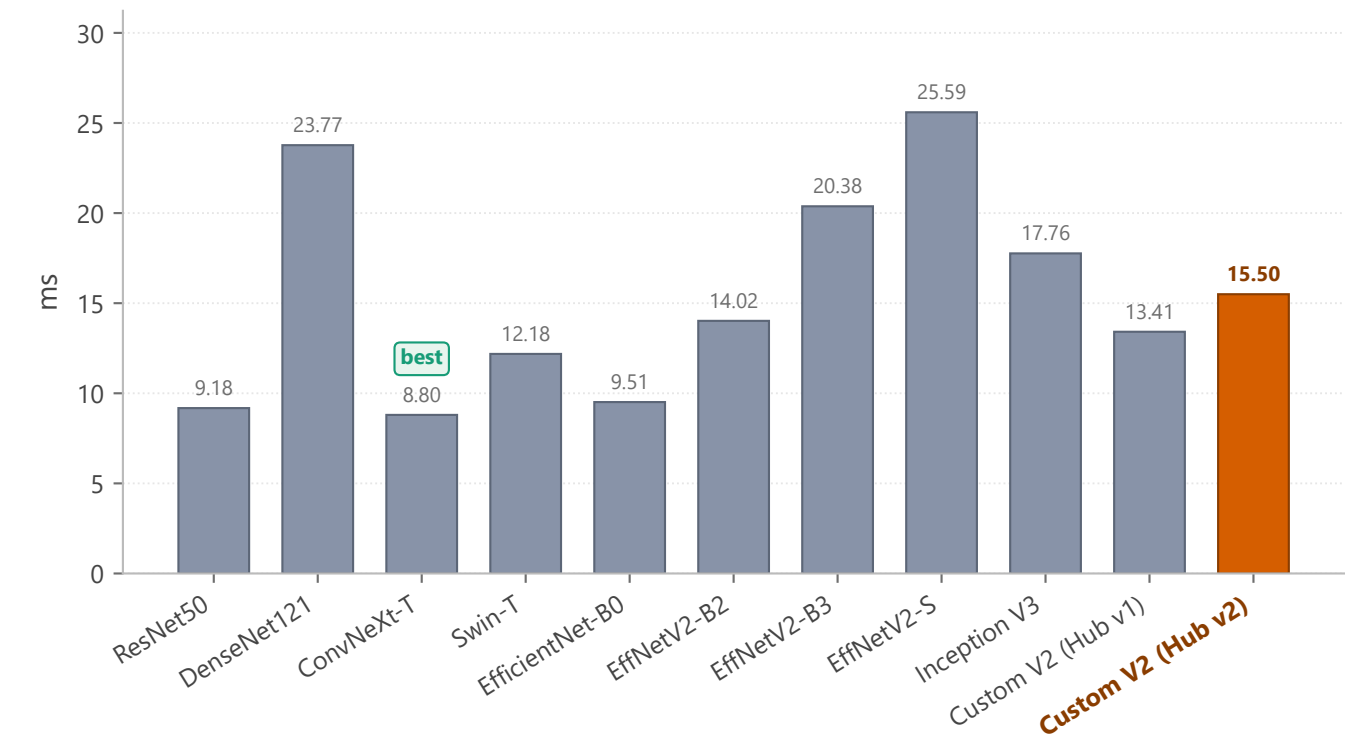
(a) GPU peak memory (lower is better)



(b) P50 inference latency (lower is better)



(c) P95 inference latency (lower is better)



(d) Total training time (lower is better)

