


AttentionHub-v2 — Triplet then SE Sequential Cascade

Channels 96 → cross-dim spatial gating (Triplet) → pure-channel gating (SE) → 112



Design rationale

- Triplet performs cross-dimensional spatial gating across the (H, W) and mixed axes.
- SE performs pure per-channel re-weighting, with no spatial role overlapping Triplet.
- Disjoint roles avoid the 98.36 % accuracy ceiling seen in BAM+Triplet and v2-EMA pairings.
- The sequential cascade preserves information — no channel concat, no LayerScale.

Result — best subtype accuracy in the paper

99.51 %

subtype accuracy 99.06 % binary accuracy

4.79 M params 0.493 GFLOPs 52.3 MB GPU 9.86 ms P50